

What It's Costing You

**The question no job board
has ever answered for you.**

Today your system reads across every listing it has collected and tells you which single skill is costing you the most interviews — in your market, for your role, weighted by your own profile. It's the smallest code in the month and the reason the project exists.

Where we left off

- 1 The model **reads** — facts out of a job description. Your code does the arithmetic.
- 2 Scores that are **deterministic, explainable, tested** — same listing twice, same number twice.
- 3 And a reason string on every row ending in **“gap: kubernetes, spark”**.

THAT FRAGMENT AT THE END

You've been generating it for two days without using it. **Today is built entirely out of it.**

Scores clustered or a bit broken? **Come anyway** — gaps come from extraction, not scoring. Today often reveals your config was out of date.

Has any job site ever told you which skill you're missing most?

Not “trending skills.” Not a course recommendation. Across every role you'd actually be a good fit for, in your city, at your level — what is the one thing that keeps disqualifying you?

WHY NOT

One listing wanting Kubernetes is **noise**. A hundred and eighty of them wanting it is the most expensive fact about your career this quarter.

One listing vs. two hundred listings

ONE LISTING ASKS

- 1 “Am I a fit for **this**?”
- 2 Per-item. Bounded. Useful.
- 3 Your Scorer already does it

TWO HUNDRED ASK

- 1 “What is the market paying for, and **where am I short**?”
- 2 Aggregate. Unavailable from any one row.
- 3 Nobody builds this screen

WHY A HUMAN STRUCTURALLY CANNOT DO THIS

To know 180 listings want Kubernetes you must read 180 listings, remember them all, and count. **Your loop did it while you were asleep.**

The most valuable output is the cheapest code.

It's a **GROUP BY**. It's counting. There is no AI in the core of it at all.

Three weeks of fetching, dedup, source layers, extraction, normalisation and scoring were not the insight. All of it was building a table where the obvious question became askable.

WHICH IS WHY ONE CLEVER PROMPT CAN'T DO THIS

A model has no idea what's in **your** market. The answer isn't in the model — it's in your data. You had to go and get the data first.

What you built today

- 1 **Canonicalisation** with an alias map you own, and an audit command to maintain it.
- 2 A Gap Analyzer ranked by **opportunity cost**, not raw frequency — so it's about you, not the market.
- 3 Every number **drillable** to the listing IDs that produced it.
- 4 Timestamped **snapshots**, so trend over time comes free.
- 5 A model that **suggests** aliases and never decides one.

Lines of AI in the analysis itself: **zero**. The AI was upstream, three days ago, reading job descriptions.

A gap report with a top ten

Computed from **50+ scored listings**. Tuesday's scoring plus today's analysis — this completes week two.

The build

- 1 Top 10 gaps, ranked by opportunity cost
- 2 **3+ aliases you** added after --audit
- 3 Every gap traceable to listing IDs
- 4 Snapshots appending, never overwriting

The different one this week

- 1 Write your **#1 gap** in your README
- 2 ...and what you're going to **do** about it
- 3 One line. Not graded on the plan.
- 4 A report you don't act on is a report you didn't need.

Rank it two ways. Point at what moved.

Same data, two orderings — raw frequency versus opportunity cost. **Only one of them is about you.**

RULE 24, PROVEN ON YOUR OWN DATA

BY RAW COUNT

1 terraform 14

2 kubernetes 31

...

6 airflow 19

BY OPPORTUNITY COST

1 kubernetes cost 24.1 mean 79

2 airflow cost 14.8 mean 74

6 terraform cost 3.2 mean 28 ← moved

Terraform appears in 14 listings — and they're ones you score in the twenties on. It's a real pattern in the market and it is not **your** problem. If your two orderings come out identical, say so — that's an honest result.

BADGE UNLOCKED



The Decoder

For the people who stopped guessing what to learn next. Your system reads the market in aggregate and tells you what it's costing you — with a number, and with receipts.

Two down. **The Tracker** ✓ · **The Decoder** ✓ · **The Automator** · **The Edge** All four plus a verified public ship → certificate.

Bring something broken.

Live debugging — **[day, date · time]**. Same format: you bring a build that doesn't work, I debug it on screen.

THIS WEEK ESPECIALLY

Bring a gap report that looks **wrong**. A broken aggregate is quiet — it doesn't throw, it just lies.

If your top gap is bizarre or your counts look impossible, that's the most instructive bug in the month. Diagnosing one on screen together beats three more videos from me.

No prep. **Bring the error message, or the number that looks wrong.**

DO THIS IN THE NEXT HOUR

Two things while it's fresh

1

Run **--audit** and add three aliases.

Ten minutes, and it will change your headline number. Go and look at your real strings — you have duplicates in there right now.

2

Post your gap report — top five, with the counts. **#MyEdge**

And say what it told you. Everyone reading it has the same question about their own career and no way to answer it.

POST IT EVEN IF IT'S UNCOMFORTABLE

“Turns out I'm missing the one thing thirty jobs want” is a far better post than a green checkmark — and it's the kind of honesty people remember.

Everything you saw today was computed correctly.

The arithmetic is right. There are tests. I can prove every number. And it was all computed from facts an LLM extracted three days ago.

SO HERE'S THE FAILURE THAT HAS NO ERROR MESSAGE

If your extractor missed “k8s experience preferred” on forty listings, your gap report is **confidently, precisely, drillably wrong** — and every number in it is arithmetically correct.

No exception. No error. The table prints beautifully. And you'd go learn the wrong thing — on the strength of a number you trusted **because you built it.**

Four sessions ago I said it would tell you what to learn.

Today it did. From your data, about your market, with receipts you can check. Nobody handed you that — there is no website that has it, because it's specific to you, and no product can afford to build something specific to one person.

THE ACTUAL ARGUMENT FOR BUILDING THIS

Not that agents are trendy. That there are questions worth answering that **nobody will ever answer for you** — and now you can answer them yourself.

Your system has an answer now.

Next week: an agent whose only job is to distrust the rest of your system — to look at the Scorer's output and the Gap Analyzer's output and ask “is this plausible?” before it reaches your dashboard.

But take today's win first. It's a real one.

Gap report by Sunday. Add three aliases. Come to the live session with something broken.